

```
In [32]: import numpy as np  
import pandas as pd
```

```
In [33]: # read the dataset  
df = pd.read_csv('pandas_practice_dataset.csv')
```

```
In [34]: # show the first 5 rows  
df.head()
```

Out[34]:

	id	name	age	department	salary
0	1	Ali	25	IT	5000
1	2	Sara	30	HR	6000
2	3	John	22	IT	4500
3	4	Mona	28	Finance	7000
4	5	Omar	35	IT	8000

```
In [35]: # show the last 5 rows  
df.tail()
```

Out[35]:

	id	name	age	department	salary
15	16	Nada	32	IT	6200
16	17	Karim	27	HR	5600
17	18	Huda	36	Sales	5900
18	19	Samir	34	IT	6100
19	20	Rania	28	Finance	7500

```
In [36]: # show random 5 rows
df.sample(5)
```

Out[36]:		id	name	age	department	salary
	13	14	Salma	21	Sales	4700
	3	4	Mona	28	Finance	7000
	7	8	Nour	29	Finance	7200
	16	17	Karim	27	HR	5600
	0	1	Ali	25	IT	5000

```
In [37]: # get the shape of the dataset shape = (rows, columns)
df.shape
```

Out[37]: (20, 5)

```
In [38]: # show the column names
df.columns
```

Out[38]: Index(['id', 'name', 'age', 'department', 'salary'], dtype='object')

```
In [39]: # show the row names
df.index
```

Out[39]: RangeIndex(start=0, stop=20, step=1)

```
In [40]: # See data types + null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20 entries, 0 to 19
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          20 non-null    int64  
 1   name         20 non-null    object  
 2   age          20 non-null    int64  
 3   department   20 non-null    object  
 4   salary       20 non-null    int64  
dtypes: int64(3), object(2)
memory usage: 932.0+ bytes
```

```
In [41]: # Select ONE column  
df['name'].head()
```

```
Out[41]: 0      Ali  
1      Sara  
2      John  
3      Mona  
4      Omar  
Name: name, dtype: object
```

```
In [42]: # select multiple column  
df[['name", "salary"]].head()
```

Out[42]:

	name	salary
0	Ali	5000
1	Sara	6000
2	John	4500
3	Mona	7000
4	Omar	8000

```
In [43]: # Select row by index  
df.iloc[0]
```

```
Out[43]: id          1  
         name        Ali  
         age          25  
         department    IT  
         salary       5000  
Name: 0, dtype: object
```

```
In [44]: # Select specific row & column  
df.loc[0, "name"]
```

```
Out[44]: 'Ali'
```

```
In [45]: # Employees older than 30  
df[df["age"] > 30]
```

Out[45]:		id	name	age	department	salary
	4	5	Omar	35	IT	8000
	8	9	Khaled	31	Sales	5100
	10	11	Hassan	40	Management	9000
	12	13	Tarek	33	Finance	6800
	14	15	Mostafa	38	Management	8500
	15	16	Nada	32	IT	6200
	17	18	Huda	36	Sales	5900
	18	19	Samir	34	IT	6100

```
In [46]: # Employees in IT department  
df[df["department"] == "IT"]
```

Out[46]:		id	name	age	department	salary
	0	1	Ali	25	IT	5000
	2	3	John	22	IT	4500
	4	5	Omar	35	IT	8000
	9	10	Dina	26	IT	5300
	15	16	Nada	32	IT	6200
	18	19	Samir	34	IT	6100

```
In [47]: # Multiple conditions  
df[(df["age"] > 25) & (df["salary"] > 6000)]
```

Out[47]:		id	name	age	department	salary
	3	4	Mona	28	Finance	7000
	4	5	Omar	35	IT	8000
	7	8	Nour	29	Finance	7200
	10	11	Hassan	40	Management	9000
	12	13	Tarek	33	Finance	6800
	14	15	Mostafa	38	Management	8500
	15	16	Nada	32	IT	6200
	18	19	Samir	34	IT	6100
	19	20	Rania	28	Finance	7500

```
In [48]: # Sort by salary  
df.sort_values("salary")
```

Out[48]:		id	name	age	department	salary
	2	3	John	22	IT	4500
	11	12	Yara	23	HR	4600
	13	14	Salma	21	Sales	4700
	6	7	Ahmed	24	Sales	4800
	0	1	Ali	25	IT	5000
	8	9	Khaled	31	Sales	5100
	9	10	Dina	26	IT	5300
	5	6	Lina	27	HR	5500
	16	17	Karim	27	HR	5600
	17	18	Huda	36	Sales	5900
	1	2	Sara	30	HR	6000
	18	19	Samir	34	IT	6100
	15	16	Nada	32	IT	6200
	12	13	Tarek	33	Finance	6800
	3	4	Mona	28	Finance	7000
	7	8	Nour	29	Finance	7200
	19	20	Rania	28	Finance	7500
	4	5	Omar	35	IT	8000
	14	15	Mostafa	38	Management	8500
	10	11	Hassan	40	Management	9000

```
In [49]: # Sort descending
df.sort_values("salary", ascending=False)
```

Out[49]:		id	name	age	department	salary
	10	11	Hassan	40	Management	9000
	14	15	Mostafa	38	Management	8500
	4	5	Omar	35	IT	8000
	19	20	Rania	28	Finance	7500
	7	8	Nour	29	Finance	7200
	3	4	Mona	28	Finance	7000
	12	13	Tarek	33	Finance	6800
	15	16	Nada	32	IT	6200
	18	19	Samir	34	IT	6100
	1	2	Sara	30	HR	6000
	17	18	Huda	36	Sales	5900
	16	17	Karim	27	HR	5600
	5	6	Lina	27	HR	5500
	9	10	Dina	26	IT	5300
	8	9	Khaled	31	Sales	5100
	0	1	Ali	25	IT	5000
	6	7	Ahmed	24	Sales	4800
	13	14	Salma	21	Sales	4700
	11	12	Yara	23	HR	4600
	2	3	John	22	IT	4500

```
In [50]: # Average salary
df["salary"].mean()
```

out[50]: np.float64(6165.0)

```
In [51]: # Maximum salary  
df["salary"].max()
```

Out[51]: 9000

```
In [52]: # Minimum salary  
df["salary"].min()
```

Out[52]: 4500

```
In [56]: # Add new column  
df["bonus"] = df["salary"] * 0.10  
df.head()
```

Out[56]:

	id	name	age	department	salary	bonus
0	1	Ali	25	IT	5000	500.0
1	2	Sara	30	HR	6000	600.0
2	3	John	22	IT	4500	450.0
3	4	Mona	28	Finance	7000	700.0
4	5	Omar	35	IT	8000	800.0

```
In [57]: # Increase age by 1 year  
df["age"] = df["age"] + 1  
df.head()
```

Out[57]:

	id	name	age	department	salary	bonus
0	1	Ali	26	IT	5000	500.0
1	2	Sara	31	HR	6000	600.0
2	3	John	23	IT	4500	450.0
3	4	Mona	29	Finance	7000	700.0
4	5	Omar	36	IT	8000	800.0

```
In [58]: # Average salary per department
df.groupby("department")["salary"].mean()
```

```
Out[58]: department
          Finance    7125.0
          HR        5425.0
          IT        5850.0
          Management 8750.0
          Sales      5125.0
Name: salary, dtype: float64
```

```
In [59]: # Count employees per department
df.groupby("department")["id"].count()
```

```
Out[59]: department
          Finance    4
          HR        4
          IT        6
          Management 2
          Sales      4
Name: id, dtype: int64
```

```
In [60]: # Save to new CSV
df.to_csv("output.csv", index=False)
```

```
In [61]: # Show employees from IT department
df[df["department"] == "IT"]
```

	id	name	age	department	salary	bonus
0	1	Ali	26	IT	5000	500.0
2	3	John	23	IT	4500	450.0
4	5	Omar	36	IT	8000	800.0
9	10	Dina	27	IT	5300	530.0
15	16	Nada	33	IT	6200	620.0
18	19	Samir	35	IT	6100	610.0

```
In [62]: # Find the highest salary
df[df["salary"] == df["salary"].max()]
```

	id	name	age	department	salary	bonus
10	11	Hassan	41	Management	9000	900.0

```
In [63]: # Find employees older than 30  
df[df["age"] > 30]
```

Out[63]:		id	name	age	department	salary	bonus
	1	2	Sara	31	HR	6000	600.0
	4	5	Omar	36	IT	8000	800.0
	8	9	Khaled	32	Sales	5100	510.0
	10	11	Hassan	41	Management	9000	900.0
	12	13	Tarek	34	Finance	6800	680.0
	14	15	Mostafa	39	Management	8500	850.0
	15	16	Nada	33	IT	6200	620.0
	17	18	Huda	37	Sales	5900	590.0
	18	19	Samir	35	IT	6100	610.0

```
In [64]: # Calculate average salary  
df["salary"].mean()
```

```
Out[64]: np.float64(6165.0)
```

```
In [65]: # Sort employees by age (descending)
df.sort_values("age", ascending=False)
```

Out[65]:		id	name	age	department	salary	bonus
	10	11	Hassan	41	Management	9000	900.0
	14	15	Mostafa	39	Management	8500	850.0
	17	18	Huda	37	Sales	5900	590.0
	4	5	Omar	36	IT	8000	800.0
	18	19	Samir	35	IT	6100	610.0
	12	13	Tarek	34	Finance	6800	680.0
	15	16	Nada	33	IT	6200	620.0
	8	9	Khaled	32	Sales	5100	510.0
	1	2	Sara	31	HR	6000	600.0
	7	8	Nour	30	Finance	7200	720.0
	3	4	Mona	29	Finance	7000	700.0
	19	20	Rania	29	Finance	7500	750.0
	16	17	Karim	28	HR	5600	560.0
	5	6	Lina	28	HR	5500	550.0
	9	10	Dina	27	IT	5300	530.0
	0	1	Ali	26	IT	5000	500.0
	6	7	Ahmed	25	Sales	4800	480.0
	11	12	Yara	24	HR	4600	460.0
	2	3	John	23	IT	4500	450.0
	13	14	Salma	22	Sales	4700	470.0

```
In [66]: # Highest salary per department
df.groupby("department")["salary"].max()
```

Out[66]: department

Finance	7500
HR	6000
IT	8000
Management	9000
Sales	5900

Name: salary, dtype: int64

```
In [67]: # Count employees in each department  
df["department"].value_counts()
```

```
Out[67]: department  
IT           6  
HR           4  
Finance      4  
Sales         4  
Management   2  
Name: count, dtype: int64
```

```
In [68]: # Show employees earning above average  
df[df["salary"] > df["salary"].mean()]
```

		id	name	age	department	salary	bonus
	3	4	Mona	29	Finance	7000	700.0
	4	5	Omar	36	IT	8000	800.0
	7	8	Nour	30	Finance	7200	720.0
	10	11	Hassan	41	Management	9000	900.0
	12	13	Tarek	34	Finance	6800	680.0
	14	15	Mostafa	39	Management	8500	850.0
	15	16	Nada	33	IT	6200	620.0
	19	20	Rania	29	Finance	7500	750.0

```
In [70]: # Check if missing exists, this will show True/False values
df.isnull()
```

Out[70]:		id	name	age	department	salary	bonus
	0	False	False	False	False	False	False
	1	False	False	False	False	False	False
	2	False	False	False	False	False	False
	3	False	False	False	False	False	False
	4	False	False	False	False	False	False
	5	False	False	False	False	False	False
	6	False	False	False	False	False	False
	7	False	False	False	False	False	False
	8	False	False	False	False	False	False
	9	False	False	False	False	False	False
	10	False	False	False	False	False	False
	11	False	False	False	False	False	False
	12	False	False	False	False	False	False
	13	False	False	False	False	False	False
	14	False	False	False	False	False	False
	15	False	False	False	False	False	False
	16	False	False	False	False	False	False
	17	False	False	False	False	False	False
	18	False	False	False	False	False	False
	19	False	False	False	False	False	False

```
In [71]: # Count missing values per column
df.isnull().sum()
```

Out[71]:

id	0
name	0
age	0
department	0
salary	0
bonus	0
dtype: int64	

```
In [72]: # Drop rows with missing data  
df.dropna()
```

Out[72]:		id	name	age	department	salary	bonus
	0	1	Ali	26	IT	5000	500.0
	1	2	Sara	31	HR	6000	600.0
	2	3	John	23	IT	4500	450.0
	3	4	Mona	29	Finance	7000	700.0
	4	5	Omar	36	IT	8000	800.0
	5	6	Lina	28	HR	5500	550.0
	6	7	Ahmed	25	Sales	4800	480.0
	7	8	Nour	30	Finance	7200	720.0
	8	9	Khaled	32	Sales	5100	510.0
	9	10	Dina	27	IT	5300	530.0
	10	11	Hassan	41	Management	9000	900.0
	11	12	Yara	24	HR	4600	460.0
	12	13	Tarek	34	Finance	6800	680.0
	13	14	Salma	22	Sales	4700	470.0
	14	15	Mostafa	39	Management	8500	850.0
	15	16	Nada	33	IT	6200	620.0
	16	17	Karim	28	HR	5600	560.0
	17	18	Huda	37	Sales	5900	590.0
	18	19	Samir	35	IT	6100	610.0
	19	20	Rania	29	Finance	7500	750.0

```
In [73]: # Drop only if specific column is missing  
df.dropna(subset=["salary"])
```

Out[73]:		id	name	age	department	salary	bonus
	0	1	Ali	26	IT	5000	500.0
	1	2	Sara	31	HR	6000	600.0
	2	3	John	23	IT	4500	450.0
	3	4	Mona	29	Finance	7000	700.0
	4	5	Omar	36	IT	8000	800.0
	5	6	Lina	28	HR	5500	550.0
	6	7	Ahmed	25	Sales	4800	480.0
	7	8	Nour	30	Finance	7200	720.0
	8	9	Khaled	32	Sales	5100	510.0
	9	10	Dina	27	IT	5300	530.0
	10	11	Hassan	41	Management	9000	900.0
	11	12	Yara	24	HR	4600	460.0
	12	13	Tarek	34	Finance	6800	680.0
	13	14	Salma	22	Sales	4700	470.0
	14	15	Mostafa	39	Management	8500	850.0
	15	16	Nada	33	IT	6200	620.0
	16	17	Karim	28	HR	5600	560.0
	17	18	Huda	37	Sales	5900	590.0
	18	19	Samir	35	IT	6100	610.0
	19	20	Rania	29	Finance	7500	750.0

```
In [74]: # Drop columns with missing data  
df.dropna(axis=1)
```

Out[74]:		id	name	age	department	salary	bonus
	0	1	Ali	26	IT	5000	500.0
	1	2	Sara	31	HR	6000	600.0
	2	3	John	23	IT	4500	450.0
	3	4	Mona	29	Finance	7000	700.0
	4	5	Omar	36	IT	8000	800.0
	5	6	Lina	28	HR	5500	550.0
	6	7	Ahmed	25	Sales	4800	480.0
	7	8	Nour	30	Finance	7200	720.0
	8	9	Khaled	32	Sales	5100	510.0
	9	10	Dina	27	IT	5300	530.0
	10	11	Hassan	41	Management	9000	900.0
	11	12	Yara	24	HR	4600	460.0
	12	13	Tarek	34	Finance	6800	680.0
	13	14	Salma	22	Sales	4700	470.0
	14	15	Mostafa	39	Management	8500	850.0
	15	16	Nada	33	IT	6200	620.0
	16	17	Karim	28	HR	5600	560.0
	17	18	Huda	37	Sales	5900	590.0
	18	19	Samir	35	IT	6100	610.0
	19	20	Rania	29	Finance	7500	750.0

```
In [75]: # Fill with a fixed value  
df.fillna(0)
```

Out[75]:		id	name	age	department	salary	bonus
	0	1	Ali	26	IT	5000	500.0
	1	2	Sara	31	HR	6000	600.0
	2	3	John	23	IT	4500	450.0
	3	4	Mona	29	Finance	7000	700.0
	4	5	Omar	36	IT	8000	800.0
	5	6	Lina	28	HR	5500	550.0
	6	7	Ahmed	25	Sales	4800	480.0
	7	8	Nour	30	Finance	7200	720.0
	8	9	Khaled	32	Sales	5100	510.0
	9	10	Dina	27	IT	5300	530.0
	10	11	Hassan	41	Management	9000	900.0
	11	12	Yara	24	HR	4600	460.0
	12	13	Tarek	34	Finance	6800	680.0
	13	14	Salma	22	Sales	4700	470.0
	14	15	Mostafa	39	Management	8500	850.0
	15	16	Nada	33	IT	6200	620.0
	16	17	Karim	28	HR	5600	560.0
	17	18	Huda	37	Sales	5900	590.0
	18	19	Samir	35	IT	6100	610.0
	19	20	Rania	29	Finance	7500	750.0

```
In [78]: # Fill with mean (very common)
df["salary"].fillna(df["salary"].mean())
```

```
Out[78]: 0      5000
         1      6000
         2      4500
         3      7000
         4      8000
         5      5500
         6      4800
         7      7200
         8      5100
         9      5300
        10     9000
        11     4600
        12     6800
        13     4700
        14     8500
        15     6200
        16     5600
        17     5900
        18     6100
        19     7500
Name: salary, dtype: int64
```

```
In [80]: # Fill with most frequent value
df["department"].fillna(df["department"].mode()[0])
```

```
Out[80]: 0          IT
         1          HR
         2          IT
         3      Finance
         4          IT
         5          HR
         6      Sales
         7      Finance
         8      Sales
         9          IT
        10    Management
        11          HR
        12    Finance
        13      Sales
        14    Management
        15          IT
        16          HR
        17      Sales
        18          IT
        19      Finance
Name: department, dtype: object
```

```
In [81]: # Check duplicates  
df.duplicated()
```

```
Out[81]: 0    False  
1    False  
2    False  
3    False  
4    False  
5    False  
6    False  
7    False  
8    False  
9    False  
10   False  
11   False  
12   False  
13   False  
14   False  
15   False  
16   False  
17   False  
18   False  
19   False  
dtype: bool
```

```
In [82]: # Count duplicates  
df.duplicated().sum()
```

```
Out[82]: np.int64(0)
```

```
In [83]: # Remove duplicates  
df.drop_duplicates(inplace=True)
```

```
In [84]: # Remove duplicates based on column
df.drop_duplicates(subset=["id"])
```

Out[84]:		id	name	age	department	salary	bonus
	0	1	Ali	26	IT	5000	500.0
	1	2	Sara	31	HR	6000	600.0
	2	3	John	23	IT	4500	450.0
	3	4	Mona	29	Finance	7000	700.0
	4	5	Omar	36	IT	8000	800.0
	5	6	Lina	28	HR	5500	550.0
	6	7	Ahmed	25	Sales	4800	480.0
	7	8	Nour	30	Finance	7200	720.0
	8	9	Khaled	32	Sales	5100	510.0
	9	10	Dina	27	IT	5300	530.0
	10	11	Hassan	41	Management	9000	900.0
	11	12	Yara	24	HR	4600	460.0
	12	13	Tarek	34	Finance	6800	680.0
	13	14	Salma	22	Sales	4700	470.0
	14	15	Mostafa	39	Management	8500	850.0
	15	16	Nada	33	IT	6200	620.0
	16	17	Karim	28	HR	5600	560.0
	17	18	Huda	37	Sales	5900	590.0
	18	19	Samir	35	IT	6100	610.0
	19	20	Rania	29	Finance	7500	750.0

```
In [85]: # Remove extra spaces
df["name"] = df["name"].str.strip()
```

```
In [86]: # Convert to lowercase
df["department"] = df["department"].str.lower()
```

```
In [87]: # Convert to uppercase
df["department"] = df["department"].str.upper()
```

```
In [88]: # Check data types  
df.dtypes
```

```
Out[88]: id          int64  
name        object  
age          int64  
department  object  
salary       int64  
bonus        float64  
dtype: object
```

```
In [89]: # Convert to integer  
df["age"] = df["age"].astype(int)
```

```
In [91]: # Convert to datetime (VERY IMPORTANT)  
# df["date"] = pd.to_datetime(df["date"])
```

```
In [93]: # Replace values  
df["department"].replace("IT Dept", "IT")
```

```
Out[93]: 0          IT  
1          HR  
2          IT  
3      FINANCE  
4          IT  
5          HR  
6      SALES  
7      FINANCE  
8      SALES  
9          IT  
10     MANAGEMENT  
11        HR  
12     FINANCE  
13      SALES  
14     MANAGEMENT  
15        IT  
16        HR  
17      SALES  
18        IT  
19      FINANCE  
Name: department, dtype: object
```

```
In [95]: # Replace multiple values  
df["department"].replace(["HR Dept", "Human Resources"], "HR")
```

```
Out[95]: 0          IT  
1          HR  
2          IT  
3      FINANCE  
4          IT  
5          HR  
6      SALES  
7      FINANCE  
8      SALES  
9          IT  
10     MANAGEMENT  
11          HR  
12     FINANCE  
13      SALES  
14     MANAGEMENT  
15          IT  
16          HR  
17      SALES  
18          IT  
19      FINANCE  
Name: department, dtype: object
```

```
In [97]: # Example: remove extremely high salaries  
df = df[df["salary"] < 100000]  
df.head()
```

		id	name	age	department	salary	bonus
0	1	Ali	26	IT		5000	500.0
1	2	Sara	31	HR		6000	600.0
2	3	John	23	IT		4500	450.0
3	4	Mona	29	FINANCE		7000	700.0
4	5	Omar	36	IT		8000	800.0

```
In [99]: # Save Cleaned Data  
df.to_csv("cleaned_data.csv", index=False)
```

```
In # Check how many missing values exist  
[100]: df.isnull().sum()
```

```
Out[100]: id      0  
          name    0  
          age     0  
          department 0  
          salary   0  
          bonus    0  
          dtype: int64
```

```
In # Fill missing salaries with the average salary  
[102]: df["salary"].fillna(df["salary"].mean())
```

```
Out[102]: 0      5000  
          1      6000  
          2      4500  
          3      7000  
          4      8000  
          5      5500  
          6      4800  
          7      7200  
          8      5100  
          9      5300  
          10     9000  
          11     4600  
          12     6800  
          13     4700  
          14     8500  
          15     6200  
          16     5600  
          17     5900  
          18     6100  
          19     7500  
          Name: salary, dtype: int64
```

```
In # Remove duplicate rows  
[103]: df.drop_duplicates(inplace=True)
```

```
In # Convert department names to uppercase  
[104]: df["department"] = df["department"].str.upper()
```

```
In # Remove extra spaces from names  
[105]: df["name"] = df["name"].str.strip()
```

```
In # Full Cleaning Pipeline (REAL ETL STYLE)
[107]: df = pd.read_csv("pandas_practice_dataset.csv")

# Check data
df.info()

# Handle missing values
df["salary"].fillna(df["salary"].mean())

# Remove duplicates
df.drop_duplicates(inplace=True)

# Clean text
df["name"] = df["name"].str.strip()
df["department"] = df["department"].str.upper()

# Save cleaned data
df.to_csv("cleaned_data.csv", index=False)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20 entries, 0 to 19
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          20 non-null    int64  
 1   name         20 non-null    object  
 2   age          20 non-null    int64  
 3   department   20 non-null    object  
 4   salary       20 non-null    int64  
 dtypes: int64(3), object(2)
 memory usage: 932.0+ bytes
```

In []: